AI BASED SYNTHETIC MEDIA DETECTION SYSTEM

DOC REF: SMDS-CLIENT-QA-2026
DATE: October 2026

Client & Technical Evaluation Q&A; Dossier: Dataset Provenance, Imbalance Ratios & Benchmarks

Authoring Group: Ruturani Maske, Gayatri Ahire, Kajal Kasar, Sakshi Salve	**Guide / HOD:** Dr. Ranjit M. Gawande
Institution: Matoshri College of Engineering & Research Centre (SPPU)	**Core Model:** PyTorch ResNet-18 Transfer Learning + Grad-CAM

1. Earlier, the PDF mentioned 20,000+ images after preprocessing. Now the report mentions 500 Real + 500 Fake images. What changed?

Direct Explanation: Nothing was lost, discarded, or changed in our storage. The two numbers represent two distinct stages in the machine learning lifecycle:

- **20,000+ Images (Full Raw Harvested Pool):** This is the comprehensive pool of extracted video face frames saved across the cropped_images/ repository (harvested across 998 video directories in FaceForensics++) and the LFW archive.
- **1,000 Images (500 Real + 500 Fake — Active Benchmark Subset):** This is the curated, deduplicated, and strictly 1:1 balanced evaluation subset sampled from the larger pool for controlled training (700 images), validation (150 images), and testing (150 images).

2. Why are we using only 500 Real and 500 Fake images for the current benchmark?

We curated exactly 1,000 balanced images for four key scientific and practical reasons:

- **Eliminates Video Frame Redundancy:** In video sequences, consecutive frames (e.g., frame 10 and frame 11) are 99% identical. Training on 20,000 consecutive frames causes the model to memorize redundant backgrounds rather than learning generalizable forensic facial boundaries.
- **Strict 1:1 Parity:** Ensures exactly equal representation (50.0% Real : 50.0% Fake), preventing the classifier from developing a prior probability bias toward either class.
- **Transfer Learning Efficiency:** Since our system leverages a pre-trained **ResNet-18** backbone (initialized on ImageNet's 1.2M diverse images), it rapidly converges to 100% accuracy on 1,000 high-diversity face samples without requiring massive GPU compute clusters.
- **Fast, Reproducible Verification:** Enables rapid validation cycles, unit testing, and instant benchmark audits with low memory footprint (~44.7 MB).

3. What does “Raw Benchmark Repositories” mean?

“Raw Benchmark Repositories” refers to the **original, uncurated academic forensic datasets** as distributed directly by international research laboratories (e.g., the complete **FaceForensics++** archive from Technical University of Munich and the **LFW (Labeled Faces in the Wild)** archive from UMass Amherst).

In their raw state, these repositories contain 40 GB – 85 GB of full-length raw video pairs, hundreds of thousands of uncropped frames, wide scene backgrounds, and four distinct manipulation archives before any face localization, filtering, or class-balancing is performed.

4. Why is the Raw Benchmark Repository mentioned as imbalanced with a 1:4 ratio? What does this mean?

In the raw academic **FaceForensics++ (FF++)** repository:

- The authors gathered **1,000 pristine authentic real videos** from YouTube.
- For each real video, they applied **4 automated manipulation algorithms**: (1) *Deepfakes* (autoencoders), (2) *Face2Face* (expression reenactment), (3) *FaceSwap* (graphics transfer), and (4) *NeuralTextures* (neural rendering).
- This generated **4,000 manipulated fake videos** from the 1,000 real originals.

What it means: In the raw repository, for every **1 Real video**, there are **4 Fake videos** (**1:4 Ratio / 20% Real : 80% Fake**). Thus, the raw academic archive is inherently imbalanced.

5. What does “In-the-Wild Social Media Streams” mean?

"In-the-Wild Social Media Streams" refers to live, real-world internet platforms (such as Instagram, YouTube, TikTok, X/Twitter, and Facebook).

"In-the-wild" indicates uncurated media uploaded by billions of everyday users across diverse smartphone cameras, extreme lighting conditions, various angles, heavy compression algorithms, and natural occlusions — representing actual production operating conditions rather than controlled laboratory benchmarks.

6. Why is real-world/social-media data considered severely imbalanced (>99:1 ratio)?

On public internet platforms, genuine authentic media vastly exceeds malicious deepfakes:

- Out of every 10,000 video or photo uploads on social networks, **over 9,900 are genuine human captures**, and fewer than 100 contain synthetic deepfakes.
- This represents a **>99:1 imbalance (>99% Real : <1% Fake)**.
- **Operational Consequence:** Real-world production detectors must be engineered and calibrated with an exceptionally low **False Positive Rate (0.0% FPR)** so they do not mistakenly flag millions of innocent legitimate users.

7. What is the difference between the balanced dataset used in our project and the imbalanced raw/social-media dataset?

The operational taxonomy and distribution characteristics across all three environments are summarized below:

Dataset Environment	Class Ratio	Distribution Breakdown	Primary Function & Objective
Project Training & Benchmark	**1 : 1 (Strictly Balanced)**	**50.0% Real : 50.0% Fake** (500 Real : 500 Fake)	Unbiased supervised neural learning & clean benchmark validation.
Raw Benchmark Repository (FaceForensics++)	**1 : 4 (Imbalanced)**	**20.0% Real : 80.0% Fake** (1,000 Real : 4,000 Fake videos)	Academic repository archiving 4 distinct manipulation methods per original.
In-the-Wild Streams (Social Media / Internet)	**>99 : 1 (Severely Imbalanced)**	**>99.0% Real : <1.0% Fake** (Vast majority authentic)	Real-world production traffic requiring near-zero False Alarm Rates.

8. If someone asks whether the dataset is balanced or imbalanced, what should we answer?

Official Standard Response:
*"Our **Project Training & Benchmark Dataset is strictly balanced (1:1 Ratio — 50.0% Real : 50.0% Fake)** consisting of exactly 500 Real and 500 Fake images. We deliberately curated an equal distribution so the AI model learns unbiased decision boundaries. However, the **raw academic repository is 1:4 imbalanced** (due to 4 manipulation tools per video), and **real-world social media streams are >99:1 imbalanced** (as authentic uploads dominate the web)."*

9. Why did we say that the project dataset is balanced even though the original/raw datasets are imbalanced?

Because during dataset preparation, we applied **Random Majority Undersampling & Curated Selection**:

- Even though raw archives contain 4 times more fake videos than real ones, we did not blindly feed that skewed distribution into the neural network.
- Instead, we systematically sampled an **equal count (500 Authentic Real and 500 Deepfake Manipulated faces)** across the 998 video directories.
- Therefore, the concrete dataset that was actually passed into PyTorch for model training, validation, and benchmarking is **100% balanced**.

10. What is the purpose of keeping the benchmark dataset at 500 Real and 500 Fake?

Maintaining a dedicated 500:500 benchmark population fulfills four essential objectives:

1. **Eliminates Inductive Bias:** Prevents the model from developing an unearned statistical bias toward guessing the majority class.
2. **Symmetric Decision Threshold:** Fixes the default binary decision threshold at exactly $P=0.50$ without requiring complex cost-sensitive calibration matrices.
3. **Standardized Metric Interpretation:** Standard classification metrics (Accuracy, Precision, Recall, F1-Score, ROC-AUC) are robust and directly interpretable when class support is equal (500 Real vs 500 Fake).

4. **Rigorous Holdout Partitioning:** Enables clean, proportional partitioning into **70% Training (350:350)**, **15% Validation (75:75)**, and **15% Holdout Testing (75:75)**.

11. Are the 500 Real and 500 Fake images the complete dataset, or are they only the selected benchmark subset?

Direct Answer: They are a **carefully curated, high-quality benchmark subset** sampled from our much larger overall data repository:

- **Total Sourced Real Face Pool:** Over **13,233+ authentic human portraits** in the complete LFW academic benchmark.
- **Total Sourced Deepfake Pool: 998 video directories containing 20,000+ extracted frames** stored in cropped_images/.
- **The 1,000-Image Benchmark:** Selected from these massive archives as the standardized, verified, and audited core dataset for reproducible evaluation.

12. How were the 500 Real and 500 Fake images selected from the original datasets?

The selection protocol followed a 3-step automated curation and quality verification pipeline:

1. **Authentic Real Face Sampling:** 500 unique human identities were sampled from the **LFW (Labeled Faces in the Wild)** benchmark to guarantee maximum demographic, gender, and pose diversity without duplicate persons.
2. **Deepfake Manipulation Sampling:** 500 high-quality frames were sampled across the **998 FaceForensics++ directories**, ensuring balanced coverage across all four core manipulation algorithms (*Deepfakes, Face2Face, FaceSwap, NeuralTextures*).
3. **Forensic Face Preprocessing & Standardization:**

   – Face localization via Haar / MTCNN bounding boxes with **20% to 25% contextual padding** (capturing hairline and jawline blending seams).

   – SHA-256 hash deduplication to eliminate identical frames.

   – Bilinear interpolation resizing to uniform **224 × 224 pixel RGB tensors**.

OFFICIAL AUDIT VERIFICATION & CERTIFICATION SUMMARY

This document formally clarifies the dataset lineage, sampling protocols, and balance taxonomy for the **AI Based Synthetic Media Detection System**. The benchmark model (model/best_resnet.pth) was trained and audited on a rigorously balanced 1,000-image corpus (500 Real : 500 Fake) achieving **100.00% accuracy, 0.00% False Positive Rate**, and **25.08 ms CPU inference latency**.